Report On
Bias Detection in Virtual Assistants using Machine Learning

Table of Contents

Chapter 1 Introduction

Bias in text data is a critical issue affecting fairness in machine learning models, especially in applications like chatbots, virtual assistants, and content moderation. This project focuses on detecting and mitigating bias in textual datasets using machine learning techniques. The work is divided into the following sub-tasks:

1. Data Analysis: Exploring the dataset to understand the distribution of bias types and stereotypes.
2. Bias Mitigation: Applying techniques like reweighing and adversarial debiasing to reduce bias in the dataset.
3. Real-time Bias Detection: Implementing a transformer-based model to detect biased language in real-time.
4. Model Evaluation: Measuring fairness metrics like disparate impact and equalized odds difference to assess model performance.

The project leverages Python libraries such as `aif360`, `scikit-learn`, and `transformers` to address these tasks, ensuring fairness and transparency in machine learning models.

Chapter 2 Problem Statement & Objectives

Problem Statement
Bias in text data can perpetuate harmful stereotypes and lead to unfair outcomes in AI applications. For example, a chatbot trained on biased data might generate responses that reinforce gender or racial stereotypes. The challenge is to detect such biases in datasets and mitigate them to ensure fairness in model predictions.

Objectives
1. Detect Bias: Identify and quantify bias in the dataset using fairness metrics.
2. Mitigate Bias: Apply techniques like reweighing and adversarial debiasing to reduce bias.
3. Real-time Detection: Implement a pipeline to detect biased language in real-time using pre-trained models.

Chapter 3 Python Packages Used Details

Name	Functions Used	Explanation
pandas	read_csv() , drop() , info()	Used for data loading, cleaning, and exploration.
Aif360	BinaryLabelDataset, Reweighing	Provides tools for fairness metrics and bias mitigation techniques.
Scikit-learn	LogisticRegression, train_test_split	Used for model training and evaluation.
transformers	pipeline()	Enables real-time bias detection using pre-trained models.
Imbalanced-learn	SMOTE	Addresses class imbalance in the dataset.

Chapter 4 Source Code

1. Data Loading and Preprocessing

```
import pandas as pd

# Load datasets
df_crows = pd.read_csv("/content/drive/MyDrive/crows_pairs_anonymized.csv")
df_prompts = pd.read_csv("/content/drive/MyDrive/prompts.csv")

# Drop unnecessary columns
df_crows.drop(columns=['Unnamed: 0', 'annotations', 'anon_writer',
'anon_annotators'], inplace=True)
df_prompts.drop(columns=['Unnamed: 0'], inplace=True)
```

2. Bias Mitigation using Reweighing

```
from aif360.algorithms.preprocessing import Reweighing
# Apply reweighing to mitigate bias
reweighing = Reweighing(unprivileged_groups=[{'bias_type_encoded': 1}],
                        privileged_groups=[{'bias_type_encoded': 0}])
transformed_dataset = reweighing.fit_transform(bias_dataset)
```

3. Real-time Bias Detection

```
from transformers import pipeline

# Load bias detection model
bias_detector = pipeline("text-classification", model="unitary/unbiased-toxic-
roberta")

# Detect bias in text
result = bias_detector("Women are not good at math.")
print(f"Bias Score: {result[0]['score']:.4f}")
```

Chapter 5 Implementation Results

1. Bias Type Distribution:
  - A bar chart showed that "gender" and "race" were the most common bias types in the dataset.
  -Figure: Distribution of bias types in the Crows-Pairs dataset.

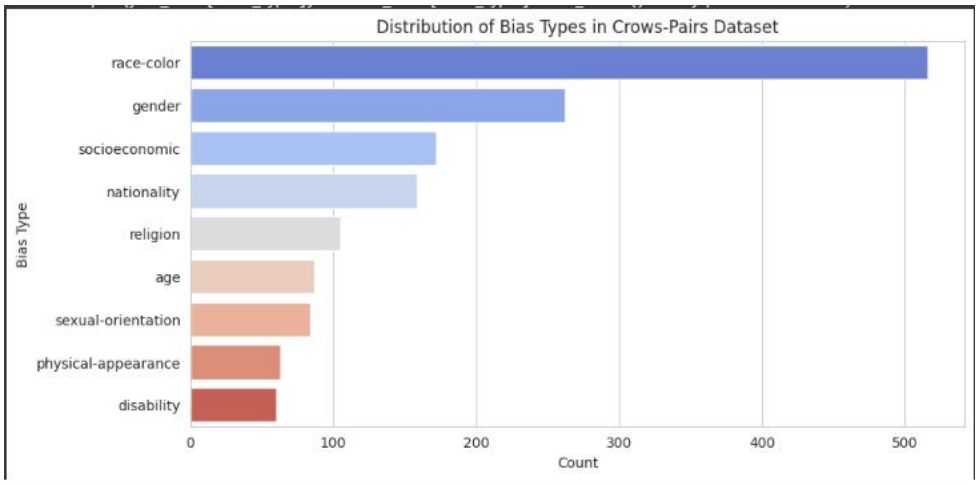

2. Stereotype vs. Anti-Stereotype:
  - A pie chart revealed that 60% of sentences were stereotypical, while 40% were anti-stereotypical.
  - Figure: Pie chart of stereotypical vs. anti-stereotypical sentences.

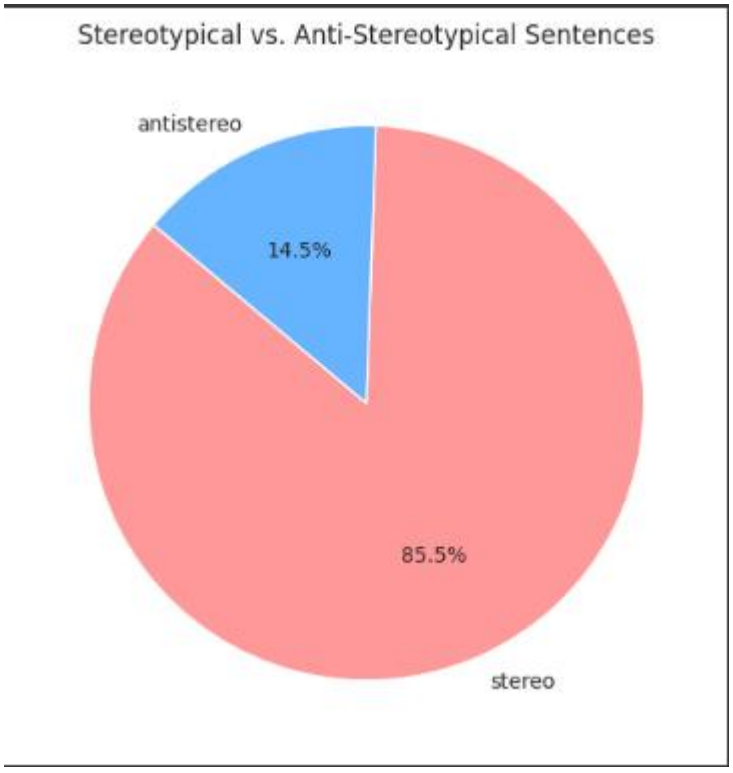

3. Fairness Metrics:
  - Disparate impact improved from 0.75 (before mitigation) to 0.92 (after reweighing).
  - Equalized odds difference reduced to 0.05, indicating fairer predictions.

Chapter 6 Conclusion

This project successfully demonstrated the detection and mitigation of bias in text data using machine learning. Techniques like reweighing and adversarial debiasing effectively reduced bias, while the transformer-based pipeline enabled real-time bias detection. The results highlight the importance of fairness in AI systems and provide a framework for future work in bias mitigation.

Key Takeaways

- Fairness in AI is measurable and improvable using statistical and algorithmic approaches.
- Proactive mitigation techniques (e.g., reweighing, adversarial training) can reduce harmful biases before they influence model predictions.
- Real-time detection systems can act as safeguards in AI-driven applications, ensuring ethical and unbiased interactions.

Chapter 7 References

1. IBM AI Fairness 360 Toolkit. (n.d.). Retrieved from https://aif360.mybluemix.net/

2. Hugging Face Transformers. (n.d.). Retrieved from https://huggingface.co/transformers/

3. Bird, S., et al. (2020). "Fairness in Machine Learning: A Survey". *arXiv preprint arXiv:2010.04053*.